

Estática II

Lista 2

Problema

Estimar três modelos (Ridge, Lasso e Elasticnet) para explicar a variável Y (price), as demais variáveis da base de dados são todas variáveis explicativas; particione a base de dados em 80% para treino e 20% para teste; e apresente os resultados:

Separando o treino

```
> set.seed(1)
>
> indices <- createDataPartition(dataset$price, p=0.8, list=F)
> treino <- dataset[indices,]
> teste <- dataset[-indices,]
```

Alterando escala das variáveis

```

> cols <- c('price', 'age', 'parea', 'tarea', 'bath',
+           'ensuit', 'garag', 'plaz', 'park', 'trans',
+           'kidca', 'school', 'health', 'bike')
>
> pre_proc_val <- preProcess(treino[,cols], method = c("center", "scale"))
>
> treino[,cols] <- predict(pre_proc_val, treino[,cols])
> teste[,cols] <- predict(pre_proc_val, teste[,cols])
>
> print("valores de treino")
[1] "valores de treino"
> summary(treino)

   price      age      parea      tarefa      bath      ensuit
Min.   :-1.3512 Min.   :-1.0471 Min.   :-2.17312 Min.   :-1.872522 Min.   :-2.5996 Min.   :-1.5890
1st Qu.:-0.7324 1st Qu.:-0.8890 1st Qu.:-0.83396 1st Qu.:-0.882164 1st Qu.:-0.8758 1st Qu.:-0.4978
Median :-0.1389 Median :-0.3354 Median :-0.04264 Median : 0.005744 Median :-0.0139 Median :-0.4978
Mean   : 0.0000 Mean   : 0.0000 Mean   : 0.000000 Mean   : 0.000000 Mean   : 0.0000 Mean   : 0.0000
3rd Qu.: 0.4155 3rd Qu.: 0.7716 3rd Qu.: 0.74108 3rd Qu.: 0.782663 3rd Qu.: 0.8480 3rd Qu.: 0.5934
Max.   : 6.1479 Max.   : 2.9857 Max.   : 2.39220 Max.   : 3.762277 Max.   : 2.5718 Max.   : 1.6845

   garag      plaz      park      trans      kidca      school
Min.   :-2.7684 Min.   :-1.6640 Min.   :-2.2628 Min.   :-2.4925 Min.   :-3.2493 Min.   :-2.06525
1st Qu.:-1.2590 1st Qu.:-0.8550 1st Qu.:-0.7801 1st Qu.:-0.7636 1st Qu.:-0.6062 1st Qu.:-0.80119
Median : 0.2504 Median :-0.1820 Median : 0.2407 Median : 0.2202 Median : 0.2238 Median :-0.01037
Mean   : 0.0000 Mean   : 0.0000 Mean   : 0.0000 Mean   : 0.0000 Mean   : 0.0000 Mean   : 0.00000
3rd Qu.: 0.2504 3rd Qu.: 0.7331 3rd Qu.: 0.8295 3rd Qu.: 0.8089 3rd Qu.: 0.7240 3rd Qu.: 0.62998
Max.   : 3.2692 Max.   : 3.2624 Max.   : 1.8369 Max.   : 1.4371 Max.   : 2.0784 Max.   : 3.47389

   health      bike      barb      balc      elev      fitg      party
Min.   :-1.7846 Min.   :-1.7295 Min.   : 0.00 Min.   : 0.0000 Min.   : 0.0000 Min.   : 0.0000 Min.   : 0.0000
1st Qu.:-0.7214 1st Qu.:-0.7519 1st Qu.: 0.00 1st Qu.: 0.0000 1st Qu.: 0.0000 1st Qu.: 0.0000 1st Qu.: 0.0000
Median :-0.2081 Median :-0.1259 Median : 0.00 Median : 0.0000 Median : 0.0000 Median : 0.0000 Median : 0.0000
Mean   : 0.0000 Mean   : 0.0000 Mean   : 0.47 Mean   : 0.4401 Mean   : 0.2972 Mean   : 0.2926 Mean   : 0.5369
3rd Qu.: 0.5379 3rd Qu.: 0.7515 3rd Qu.: 1.00 3rd Qu.: 1.0000 3rd Qu.: 1.0000 3rd Qu.: 1.0000 3rd Qu.: 1.0000
Max.   : 3.9310 Max.   : 3.5095 Max.   : 1.00 Max.   : 1.0000 Max.   : 1.0000 Max.   : 1.0000 Max.   : 1.0000

   categ
Min.   : 0.0000
1st Qu.: 1.0000
Median : 1.0000
Mean   : 0.9562
3rd Qu.: 1.0000
Max.   : 1.0000

[1] "valores de teste"
> summary(teste)

   price      age      parea      tarefa      bath      ensuit
Min.   :-1.2572 Min.   :-1.04711 Min.   :-1.68615 Min.   :-1.68470 Min.   :-2.5996 Min.   :-1.588981
1st Qu.:-0.7310 1st Qu.:-0.88896 1st Qu.:-0.84918 1st Qu.:-0.80533 1st Qu.:-0.8758 1st Qu.:-0.497814
Median :-0.1577 Median :-0.49358 Median :-0.04264 Median :-0.07963 Median :-0.0139 Median :-0.497814
Mean   :-0.0544 Mean   :-0.08712 Mean   :-0.03041 Mean   :-0.06399 Mean   :-0.1508 Mean   :-0.008318
3rd Qu.: 0.3500 3rd Qu.: 0.73209 3rd Qu.: 0.67260 3rd Qu.: 0.64606 3rd Qu.: 0.8480 3rd Qu.: 0.593354
Max.   : 2.3357 Max.   : 2.27406 Max.   : 1.90523 Max.   : 1.91816 Max.   : 1.7099 Max.   : 1.684521

   garag      plaz      park      trans      kidca      school
Min.   :-2.76839 Min.   :-1.67190 Min.   :-1.7396 Min.   :-2.19678 Min.   :-3.00038 Min.   :-1.9975
1st Qu.:-1.25899 1st Qu.:-0.77106 1st Qu.: 0.4292 1st Qu.:-0.76358 1st Qu.:-0.51831 1st Qu.:-0.7239
Median : 0.25041 Median :-0.02648 Median : 0.5308 Median : 0.06041 Median : 0.09492 Median :-0.1170
Mean   : 0.08113 Mean   : 0.07582 Mean   : 0.2440 Mean   :-0.01099 Mean   :-0.02816 Mean   :-0.1620
3rd Qu.: 0.25041 3rd Qu.: 0.92688 3rd Qu.: 0.9718 3rd Qu.: 0.77202 3rd Qu.: 0.67673 3rd Qu.: 0.2248
Max.   : 3.26920 Max.   : 2.65944 Max.   : 1.5623 Max.   : 1.43446 Max.   : 1.47962 Max.   : 2.3377

   health      bike      barb      balc      elev      fitg
Min.   :-1.5980 Min.   :-1.7051 Min.   : 0.0000 Min.   : 0.0000 Min.   : 0.0000 Min.   : 0.0000
1st Qu.:-0.7324 1st Qu.:-0.5777 1st Qu.: 0.0000 1st Qu.: 0.0000 1st Qu.: 0.0000 1st Qu.: 0.0000
Median :-0.1705 Median : 0.1026 Median : 1.0000 Median : 0.0000 Median : 0.0000 Median : 0.0000
Mean   : 0.1157 Mean   : 0.2439 Mean   : 0.5981 Mean   : 0.4766 Mean   : 0.3551 Mean   : 0.3738
3rd Qu.: 0.8431 3rd Qu.: 0.9821 3rd Qu.: 1.0000 3rd Qu.: 1.0000 3rd Qu.: 1.0000 3rd Qu.: 1.0000
Max.   : 3.9310 Max.   : 3.1196 Max.   : 1.0000 Max.   : 1.0000 Max.   : 1.0000 Max.   : 1.0000

   party      categ
Min.   : 0.0000 Min.   : 0.0000
1st Qu.: 0.0000 1st Qu.: 1.0000
Median : 1.0000 Median : 1.0000
Mean   : 0.5794 Mean   : 0.9533
3rd Qu.: 1.0000 3rd Qu.: 1.0000
Max.   : 1.0000 Max.   : 1.0000

```

Função que calcula e retorna R^2 e RMSE

```

> eval_results <- function(true, predicted, df) {
+   SSE <- sum((predicted - true)^2)
+   SST <- sum((true - mean(true))^2)
+   R_square <- 1 - SSE / SST
+   RMSE = sqrt(SSE/nrow(df))
+
+   # Model performance metrics
+   data.frame(
+     RMSE = RMSE,
+     Rsquare = R_square
+   )
+ }

```

Regressão Ridge

```

> cols_reg <- c('price', 'age', 'parea', 'tarea', 'bath',
+              'ensuit', 'garag', 'plaz', 'park', 'trans',
+              'kidca', 'school', 'health', 'bike', 'barb',
+              'balc', 'elev', 'fitg', 'party', 'categ')
>
> dummies <- dummyvars(price ~ age+parea+tarea+bath+
+                      ensuit+garag+plaz+park+trans+kidca+
+                      school+health+bike+barb+balc+elev+fitg+
+                      party+categ,
+                      data = dataset[,cols_reg])
>
> train_dummies <- predict(dummies, newdata = treino[,cols_reg])
>
> test_dummies <- predict(dummies, newdata = teste[,cols_reg])
>
> print(dim(train_dummies)); print(dim(test_dummies))
[1] 434  19
[1] 107  19

```

i. O valor ótimo do lambda para os modelos;

O melhor lamba 0.1

```

> x = as.matrix(train_dummies)
> y_train = treino$price
>
> x_test = as.matrix(test_dummies)
> y_test = teste$price
> lambdas <- 10^seq(2, -3, by = -.1)
> ridge_lamb <- cv.glmnet(x, y_train, alpha = 0,
+                         lambda = lambdas)
> best_lambda_ridge <- ridge_lamb$lambda.min
> best_lambda_ridge
[1] 0.1

```

ii. O valor do alpha para o modelo ElasticNet;

iii. Os valores dos parâmetros para os modelos;

```
> ridge_reg[["beta"]]
19 x 1 sparse Matrix of class "dgCMatrix"
      s0
age    -0.17994688
parea  0.15907495
tarea  0.21631677
bath   0.04261451
ensuit  0.18954877
garag  0.20506605
plaz   0.04714112
park   -0.05375536
trans  0.03345027
kidca  0.01621329
school -0.00131847
health -0.01054514
bike   -0.04790725
barb   -0.06878716
balc   0.15377909
elev   -0.18946965
fitg   0.23383601
party  0.06497685
categ  0.46258571
```

iv. O R^2 e RMSE dos modelos estimados;

```
> predictions_test <- predict(ridge_reg, s = best_lambda_ridge,
+                             newx = x_test)
> eval_results(y_test, predictions_test, test)
Error in nrow(df) : object 'test' not found
> predictions_test <- predict(ridge_reg, s = best_lambda_ridge,
+                             newx = x_test)
> eval_results(y_test, predictions_test, teste)
      RMSE   Rsquare
1 0.3282966 0.8487347
~ |
```

O R^2 está próximo de 1, mas não tão próximo, o que é um bom sinal pois significa que tem menos chances de o modelo estar sofrendo overfitting. O RMSE está próximo de zero, o que é bom sinal, significa que poucos erros foram cometidos.

v. Apresente os resultados de uma predição proposta por você mesmo para os modelos (valor estimado e intervalos de confiança).

Valores para a predição

```
> price <- (median(dataset$price)-pre_proc_val[["mean"]][["price"]])/pre_proc_val[["std"]][["price"]]
> age <- (median(dataset$age)-pre_proc_val[["mean"]][["age"]])/pre_proc_val[["std"]][["age"]]
> parea <- (median(dataset$parea)-pre_proc_val[["mean"]][["parea"]])/pre_proc_val[["std"]][["parea"]]
> tarea <- (median(dataset$tarea)-pre_proc_val[["mean"]][["tarea"]])/pre_proc_val[["std"]][["tarea"]]
> bath <- (median(dataset$bath)-pre_proc_val[["mean"]][["bath"]])/pre_proc_val[["std"]][["bath"]]
> ensuit <- (median(dataset$ensuit)-pre_proc_val[["mean"]][["ensuit"]])/pre_proc_val[["std"]][["ensuit"]]
> garag <- (median(dataset$garag)-pre_proc_val[["mean"]][["garag"]])/pre_proc_val[["std"]][["garag"]]
> plaz <- (median(dataset$plaz)-pre_proc_val[["mean"]][["plaz"]])/pre_proc_val[["std"]][["plaz"]]
> park <- (median(dataset$park)-pre_proc_val[["mean"]][["park"]])/pre_proc_val[["std"]][["park"]]
> trans <- (median(dataset$trans)-pre_proc_val[["mean"]][["trans"]])/pre_proc_val[["std"]][["trans"]]
> kidca <- (median(dataset$kidca)-pre_proc_val[["mean"]][["kidca"]])/pre_proc_val[["std"]][["kidca"]]
> school <- (median(dataset$school)-pre_proc_val[["mean"]][["school"]])/pre_proc_val[["std"]][["school"]]
> health <- (median(dataset$health)-pre_proc_val[["mean"]][["health"]])/pre_proc_val[["std"]][["health"]]
> bike <- (median(dataset$bike)-pre_proc_val[["mean"]][["bike"]])/pre_proc_val[["std"]][["bike"]]
> barb <- 0
> balc <- 0
> elev <- 0
> fitg <- 0
> party <- 0
> categ <- 0
>
> # Construindo matriz com dados para predição
> our_pred <- as.matrix(data.frame(age=age,
+                                parea=parea,
+                                tarea=tarea,
+                                bath=bath,
+                                ensuit=ensuit,
+                                garag=garag,
+                                plaz=plaz,
+                                park=park,
+                                trans=trans,
+                                kidca=kidca,
+                                school=school,
+                                health=health,
+                                bike=bike,
+                                barb=barb,
+                                balc=balc,
+                                elev=elev,
+                                fitg=fitg,
+                                party=party,
+                                categ=categ))
+ ,
```

Predição Ridge

```
> predict_our_ridge <- predict(ridge_reg, s = best_lambda_ridge,
+                               newx = our_pred)
> predict_our_ridge
s1
[1,] -0.5111749
```

Intervalo de confiança

```
> n <- nrow(treino)
> m <- predict_our_ridge
> s <- pre_proc_val[["std"]][["price"]]
> dam <- s/sqrt(n)
> CIlwr_ridge <- m + (qnorm(0.025))*dam
> CIupr_ridge <- m - (qnorm(0.025))*dam
>
> CIlwr_ridge
s1
[1,] -50056.16
> CIupr_ridge
s1
[1,] 50055.14
+ ,
```

Nota-se que o valor obtido assim como o intervalo de confiança tiveram grande divergência com o valor esperado.

Regressão de Lasso

```
> lambdas <- 10^seq(2, -3, by = -.1)
>
> lasso_lamb <- cv.glmnet(x, y_train, alpha = 1,
+                         lambda = lambdas,
+                         standardize = TRUE, nfolds = 5)
> best_lambda_lasso <- lasso_lamb$lambda.min
> best_lambda_lasso
[1] 0.007943282
>
> lasso_model <- glmnet(x, y_train, alpha = 1,
+                      lambda = best_lambda_lasso,
+                      standardize = TRUE)
```

Melhor lambda: 0.007943282

```
> lasso_model[["beta"]]
19 x 1 sparse Matrix of class "dgCMatrix"
      s0
age    -0.177787541
parea  0.160244253
tarea  0.239394524
bath    0.005550922
ensuit  0.215324160
garag   0.214458904
plaz    0.043167047
park   -0.056688502
trans   0.030263567
kidca   0.008624390
school  .
health -0.002502411
bike   -0.038298740
barb   -0.066286730
balc    0.147531615
elev   -0.179246307
fitg    0.240380064
party   0.043609304
categ   0.490774990
```

Testando o Modelo Lasso

```
> predictions_test <- predict(lasso_model, s = best_lambda_lasso,
+                             newx = x_test)
> eval_results(y_test, predictions_test, teste)
      RMSE  Rsquare
1 0.3313211 0.8459347
```

O R^2 está próximo de 1, mas não tão próximo, o que é um bom sinal pois significa que tem menos chances de o modelo estar sofrendo overfitting. O RMSE está menos próximo de zero do que o modelo de regressão Ridge.

Predição Lasso

```
> predict_our_lasso <- predict(lasso_model, s = best_lambda_lasso,
+                               newx = our_pred)
> predict_our_lasso
               s1
[1,] -0.5468196
>
> n <- nrow(treino)
> m <- predict_our_lasso
> s <- pre_proc_val[["std"]][["price"]]
> dam <- s/sqrt(n)
> CIlwr_lasso <- m + (qnorm(0.025))*dam
> CIupr_lasso <- m - (qnorm(0.025))*dam
>
> CIlwr_lasso
               s1
[1,] -50056.19
> CIupr_lasso
               s1
[1,] 50055.1
```

O valor obtido assim como o intervalo de confiança tiveram grande divergência com o valor esperado, mas o resultado foi melhor que o do modelo Ridge.

Regressão ElasticNet

```
> train_cont <- trainControl(method = "repeatedcv",
+                             number = 10,
+                             repeats = 5,
+                             search = "random",
+                             verboseIter = TRUE)
> |
> elastic_reg <- train(price ~ age+parea+tarea+bath+ensuit+garag+plaz+park+trans+
+                       kidca+school+health+bike+barb+balc+elev+fitg+party+categ,
+                       data = train,
+                       method = "glmnet",
+                       tuneLength = 10,
+                       trControl = train_cont)|
```



```

cur cu 0.1227790000 0.1227790000 0.1227790000 0.1227790000 0.1227790000 0.1227790000 0.1227790000 0.1227790000
bath 0.085199550 0.0838115043 0.082680736 0.081528628 0.080398280 0.079272601 0.078032303 0.076890291
ensuit 0.166680433 0.1683688326 0.169734994 0.170996122 0.172165770 0.173304335 0.174640609 0.175825659
garag 0.180041844 0.1812585355 0.182722640 0.184068392 0.185309687 0.186496073 0.187655313 0.188749907
plaz 0.001354329 0.004030818 0.006482205 0.009054286 0.011609225 0.014174762 0.016740299 0.019305936
park -0.055565035 -0.0565156797 -0.057187165 -0.057774557 -0.058291199 -0.059054286 -0.059984208 -0.060872222
trans 0.013537230 0.0157378812 0.017617838 0.019381877 0.021040782 0.022609225 0.024139608 0.025571020
kidca 0.005320989 0.0063061550 0.007231698 0.008073819 0.008845412 0.009611625 0.010495832 0.011256081
school . . . . .
health . . . . .
bike -0.015129523 -0.0169907358 -0.018405091 -0.019699220 -0.020884566 -0.022214518 -0.023779544 -0.025180713
barb . -0.0009872164 -0.005281610 -0.009280631 -0.013013639 -0.016286637 -0.019130596 -0.021757180
balc 0.037228557 0.0383785163 0.040087885 0.041681539 0.043166537 0.044508076 0.045819284 0.047013967
elev -0.012373887 -0.0181428622 -0.023340694 -0.028225190 -0.032813899 -0.037046132 -0.040896011 -0.044478504
fitg 0.090793036 0.0921112018 0.093275213 0.094313436 0.095239701 0.096020236 0.096746684 0.097415987
party 0.026017745 0.0288202637 0.031712681 0.034451752 0.037052440 0.039391265 0.041441026 0.043344947
categ 0.034892678 0.0387663576 0.042412656 0.045835375 0.049030475 0.052047172 0.054809446 0.057376585

age -0.187116035 -0.18930028 -0.19127537 -0.192861425 -0.194380844 -0.19577849 -0.197064417 -0.198246428
parea 0.137896544 0.13848367 0.13935726 0.139538289 0.140037631 0.14050514 0.140942297 0.141351943
tarea 0.240122535 0.24076487 0.24130877 0.241910495 0.242395925 0.24283208 0.243219960 0.243565225
bath 0.075721825 0.07460580 0.07351596 0.072684901 0.071735743 0.07084584 0.070004334 0.069209639
ensuit 0.176923807 0.17794520 0.17891762 0.180108258 0.181190495 0.18219072 0.183124145 0.183993789
garag 0.189717260 0.19060946 0.19134377 0.192049756 0.192655927 0.19321625 0.193730869 0.194203366
plaz 0.008734664 0.01081382 0.01268674 0.014464379 0.016086532 0.01757732 0.018948370 0.020208278
park -0.061668120 -0.06239667 -0.06325205 -0.064496546 -0.065591479 -0.06660555 -0.067544140 -0.068411874
trans 0.026859512 0.02805502 0.02898712 0.029616795 0.030189413 0.03070971 0.031182111 0.031610762
kidca 0.011973025 0.01263233 0.01345910 0.014567491 0.015578238 0.01651134 0.017374672 0.018172488
school . . . . .
health . . . . .
bike -0.026481598 -0.02767867 -0.02863423 -0.029305467 -0.029938754 -0.03051350 -0.031037217 -0.031514381
barb -0.024187723 -0.02643207 -0.02852914 -0.030502324 -0.032312638 -0.03397742 -0.035507644 -0.036913022
balc 0.048099371 0.04910191 0.05002566 0.050992737 0.051856139 0.05265161 0.053384378 0.054058762
elev -0.047835562 -0.05095542 -0.05387882 -0.056514102 -0.058979960 -0.06125669 -0.063357273 -0.065293266
fitg 0.097941254 0.09840496 0.09880648 0.099385014 0.099829692 0.10022865 0.100585294 0.100904138
party 0.045179387 0.04689775 0.04846556 0.049689719 0.050896029 0.05201353 0.053050074 0.054010264
categ 0.059791918 0.06202857 0.06413369 0.065910528 0.067605222 0.06916508 0.070599053 0.071916466

age -0.19933201 .....
parea 0.14173503 .....
tarea 0.24387260 .....
bath 0.06846135 .....
ensuit 0.18480284 .....
garag 0.19463711 .....
plaz 0.02136524 .....
park -0.06921329 .....
trans 0.03199968 .....
kidca 0.01890897 .....
school . . . . .
health 0.01024899 .....
bike -0.03194908 .....
barb -0.03820282 .....
balc 0.05467891 .....
elev -0.06707586 .....
fitg 0.10118930 .....
party 0.05489859 .....
categ 0.07312594 .....

```

```

.....suppressing 20 columns in show(); maybe adjust 'options(max.print= *, width = *)'
.....

```

Treino

```

> predictions_train <- predict(elastic_reg, x)
> eval_results(y_train, predictions_train, train)
      RMSE  Rsquare
1 0.3981597 0.841101

```

Teste

```

> predictions_test <- predict(elastic_reg, x_test)
> eval_results(y_test, predictions_test, test)
      RMSE  Rsquare
1 0.6545521 0.7042617

```

Predição

```

>
> price_pred_elastic=(predict_our_elastic*
+                      pre_proc_val[["std"]][["price"]])+
+                      pre_proc_val[["mean"]][["price"]]
> price_pred_elastic
[1] 1020881

```

```

> n <- nrow(train)
> m <- price_pred_elastic
> s <- pre_proc_val[["std"]][["price"]]
> dam <- s/sqrt(n)
> CIlwr_elastic <- m + (qnorm(0.025))*dam
> CIupr_elastic <- m - (qnorm(0.025))*dam
>
> CIlwr_elastic
[1] 974239.1
> CIupr_elastic
[1] 1067522
`|

```

Dentre os modelos testados o que obteve melhor resultado foi o lasso. Todavia em todos modelos o valor obtido assim como o intervalo de confiança tiveram grande divergência com o valor esperado.